

Customer Propensity Modeling

Wyatt Rasmussen

DSC 680

Final Project 1

Data Imports

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
In [2]: training = pd.read_csv('data/training_sample.csv')
testing = pd.read_csv('data/testing_sample.csv')
```

Cleaning and Exploring Data

```
In [3]: print('Size of training dataset:', training.shape)
print('Size of testing dataset:', testing.shape)
```

```
Size of training dataset: (455401, 25)
Size of testing dataset: (151655, 25)
```

```
In [4]: print('Null count of training dataset:', sum(training.isna().sum()))
print('Null count of testing dataset:', sum(testing.isna().sum()))
```

```
Null count of training dataset: 0
Null count of testing dataset: 0
```

```
In [5]: training.columns
```

```
Out[5]: Index(['UserID', 'basket_icon_click', 'basket_add_list', 'basket_add_detail',
'sort_by', 'image_picker', 'account_page_click', 'promo_banner_click',
'detail_wishlist_add', 'list_size_dropdown', 'closed_minibasket_click',
'checked_delivery_detail', 'checked_returns_detail', 'sign_in',
'saw_checkout', 'saw_sizecharts', 'saw_delivery', 'saw_account_upgrade',
'saw_homepage', 'device_mobile', 'device_computer', 'device_tablet',
'returning_user', 'loc_uk', 'ordered'],
dtype='object')
```

```
In [6]: testing.columns
```

```
Out[6]: Index(['UserID', 'basket_icon_click', 'basket_add_list', 'basket_add_detail',
'sort_by', 'image_picker', 'account_page_click', 'promo_banner_click',
```


```
'detail_wishlist_add', 'list_size_dropdown', 'closed_minibasket_click',
'checked_delivery_detail', 'checked_returns_detail', 'sign_in',
'saw_checkout', 'saw_sizecharts', 'saw_delivery', 'saw_account_upgrade',
'saw_homepage', 'device_mobile', 'device_computer', 'device_tablet',
'returning_user', 'loc_uk', 'ordered'],
dtype='object')
```

In [7]: `training.dtypes`

```
Out[7]: UserID                object
basket_icon_click           int64
basket_add_list             int64
basket_add_detail           int64
sort_by                    int64
image_picker               int64
account_page_click          int64
promo_banner_click          int64
detail_wishlist_add         int64
list_size_dropdown          int64
closed_minibasket_click     int64
checked_delivery_detail      int64
checked_returns_detail      int64
sign_in                    int64
saw_checkout                int64
saw_sizecharts              int64
saw_delivery                int64
saw_account_upgrade         int64
saw_homepage                int64
device_mobile               int64
device_computer             int64
device_tablet               int64
returning_user              int64
loc_uk                     int64
ordered                    int64
dtype: object
```

```
In [8]: ## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)


columns = ['basket_icon_click', 'basket_add_list', 'basket_add_detail']
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0, 450000)
    ax[i].set_title(*[col])
```



In [9]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

columns = ['sort_by', 'image_picker', 'account_page_click']
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0, 450000)
    ax[i].set_title(*[col])
```




In [10]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

columns = ['promo_banner_click', 'detail_wishlist_add', 'list_size_dropdown']
for i, col in enumerate(columns):
```


```
graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
graph.set_ylim(0,450000)
ax[i].set_title(*[col])
```



In [11]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

columns = ['closed_minibasket_click', 'checked_delivery_detail', 'checked_return
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0,450000)
    ax[i].set_title(*[col])
```




In [12]:

```
## for multiple columns
```

```
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

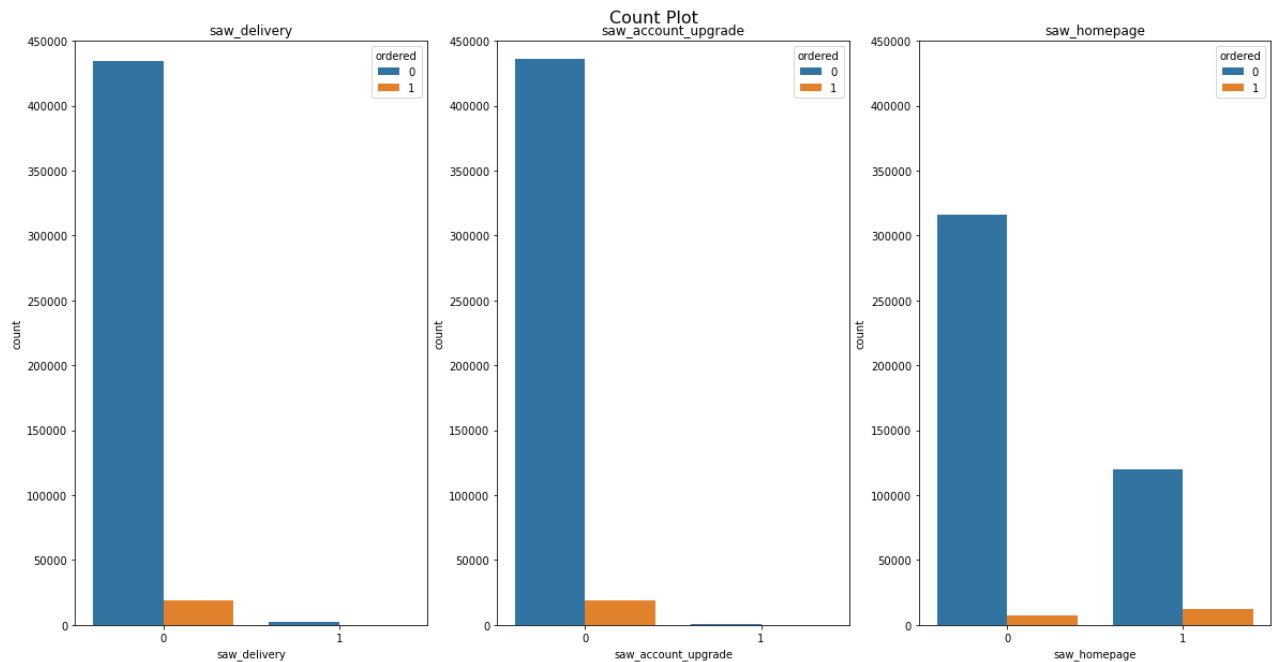
columns = ['sign_in', 'saw_checkout', 'saw_sizecharts']
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0, 450000)
    ax[i].set_title(*[col])
```



In [13]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)


columns = ['saw_delivery', 'saw_account_upgrade', 'saw_homepage']
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0, 450000)
    ax[i].set_title(*[col])
```



In [14]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

columns = ['device_mobile', 'device_computer']
for i, col in enumerate(columns):
    graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
    graph.set_ylim(0, 450000)
    ax[i].set_title(*[col])
```

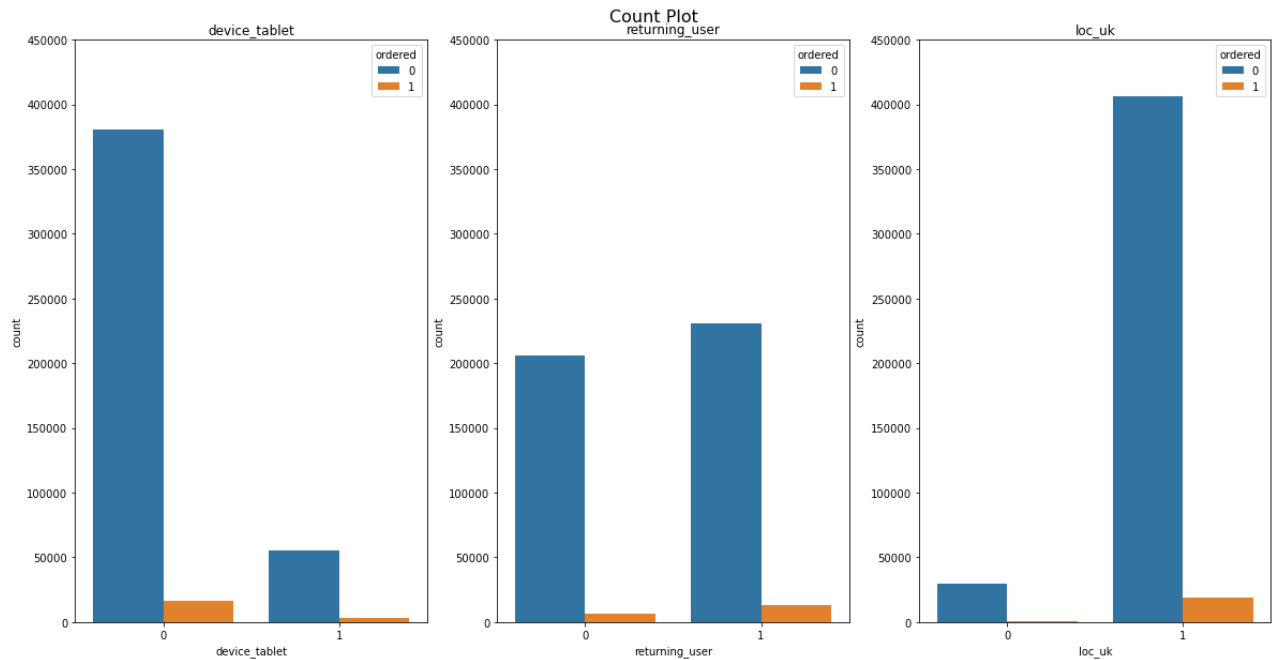


In [15]:

```
## for multiple columns
fig, ax = plt.subplots(1, 3, figsize=(20, 10))
fig.suptitle('Count Plot', fontsize=16, y=0.92)

columns = ['device_tablet', 'returning_user', 'loc_uk']
for i, col in enumerate(columns):
```

```
graph = sns.countplot(x=training[col], hue=training["ordered"], ax=ax[i])
graph.set_ylim(0,450000)
ax[i].set_title(*[col])
```



In [16]:

```
corr = training.corr()

corr.style.background_gradient(cmap='coolwarm')
```

Out[16]:

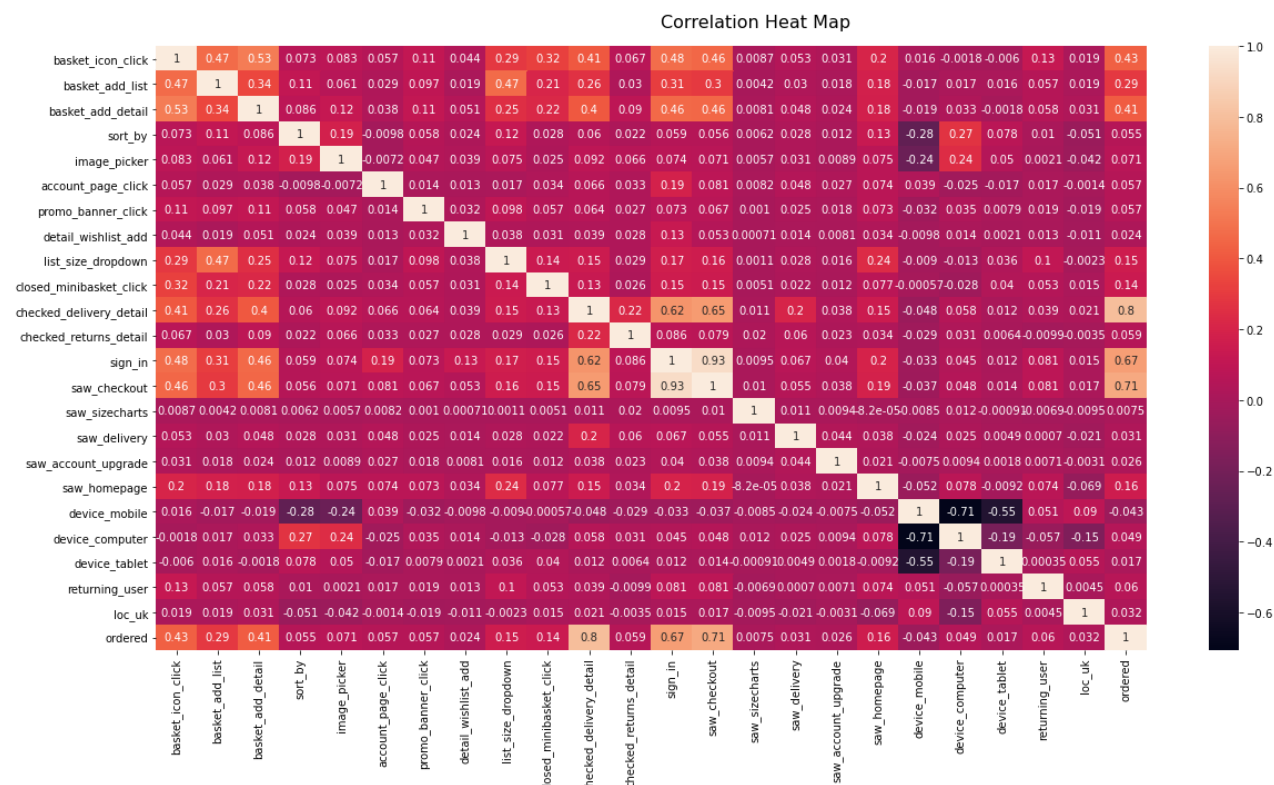
	basket_icon_click	basket_add_list	basket_add_detail	sort_by	image_picker
basket_icon_click	1.000000	0.466671	0.529947	0.073016	0.082893
basket_add_list	0.466671	1.000000	0.340968	0.106852	0.061462
basket_add_detail	0.529947	0.340968	1.000000	0.085854	0.124230
sort_by	0.073016	0.106852	0.085854	1.000000	0.185661
image_picker	0.082893	0.061462	0.124230	0.185661	1.000000
account_page_click	0.057253	0.028994	0.037502	-0.009754	-0.000000
promo_banner_click	0.109342	0.096608	0.109043	0.058155	0.000000
detail_wishlist_add	0.044153	0.019061	0.050724	0.024056	0.000000
list_size_dropdown	0.291608	0.469625	0.247205	0.124273	0.000000
closed_minibasket_click	0.323940	0.208082	0.222444	0.028453	0.000000
checked_delivery_detail	0.405787	0.264766	0.404134	0.059635	0.000000
checked_returns_detail	0.067149	0.030469	0.090434	0.022364	0.000000
sign_in	0.478834	0.312276	0.461659	0.058662	0.000000
saw_checkout	0.458774	0.297681	0.456713	0.055959	0.000000
saw_sizecharts	0.008741	0.004161	0.008101	0.006196	0.000000
saw_delivery	0.052922	0.030286	0.048410	0.028102	0.000000

	basket_icon_click	basket_add_list	basket_add_detail	sort_by	image
saw_account_upgrade	0.030764	0.018150	0.024255	0.012194	0.
saw_homepage	0.203087	0.180221	0.175138	0.128205	0.
device_mobile	0.016203	-0.017202	-0.018800	-0.278043	-0.
device_computer	-0.001757	0.016629	0.032794	0.269589	0
device_tablet	-0.006019	0.015516	-0.001799	0.078088	0.
returning_user	0.126640	0.057443	0.057680	0.010366	0.
loc_uk	0.018518	0.018797	0.030956	-0.051148	-0.
ordered	0.428334	0.287666	0.414420	0.054636	0

In [17]:

```
fig, ax = plt.subplots(1, 1, figsize=(20, 10))
fig.suptitle('Correlation Heat Map', fontsize=16, y=0.92)

sns.heatmap(training.corr(), annot = True)
plt.show()
```



In [18]:

```
training.corr()['ordered']
```

```
Out[18]: basket_icon_click    0.428334
basket_add_list    0.287666
basket_add_detail  0.414420
sort_by           0.054636
image_picker      0.071492
account_page_click 0.057279
promo_banner_click 0.056533
detail_wishlist_add 0.023516
list_size_dropdown 0.154867
```



```
closed_minibasket_click    0.140011
checked_delivery_detail    0.798720
checked_returns_detail     0.059484
sign_in                    0.665556
saw_checkout               0.708986
saw_sizecharts             0.007548
saw_delivery               0.031461
saw_account_upgrade       0.025857
saw_homepage              0.157778
device_mobile              -0.042907
device_computer            0.049208
device_tablet              0.016939
returning_user             0.060295
loc_uk                     0.031643
ordered                    1.000000
Name: ordered, dtype: float64
```

```
In [19]: training.corr()['ordered'] > 0.15
```

```
Out[19]: basket_icon_click      True
basket_add_list      True
basket_add_detail    True
sort_by              False
image_picker         False
account_page_click   False
promo_banner_click   False
detail_wishlist_add  False
list_size_dropdown   True
closed_minibasket_click False
checked_delivery_detail True
checked_returns_detail False
sign_in              True
saw_checkout         True
saw_sizecharts       False
saw_delivery         False
saw_account_upgrade  False
saw_homepage         True
device_mobile        False
device_computer      False
device_tablet        False
returning_user       False
loc_uk               False
ordered              True
Name: ordered, dtype: bool
```

```
In [20]: training.corr()['ordered'] > 0.02
```

```
Out[20]: basket_icon_click      True
basket_add_list      True
basket_add_detail    True
sort_by              True
image_picker         True
account_page_click   True
promo_banner_click   True
detail_wishlist_add  True
list_size_dropdown   True
closed_minibasket_click True
checked_delivery_detail True
checked_returns_detail True
sign_in              True
saw_checkout         True
saw_sizecharts       False
```

```
saw_delivery      True
saw_account_upgrade  True
saw_homepage      True
device_mobile     False
device_computer   True
device_tablet     False
returning_user    True
loc_uk            True
ordered           True
Name: ordered, dtype: bool
```

Feature Selection and Separating Predictors from Target Variable

Methods

For the feature selection I would like to try 2 different methods. 1st I would like to take the variables over 0.15 correlation with ordered which would be 8 features. 2nd I would like to take the variables that are over 0.02, which would be 20 features. This will change and influence the number of features that are being used in the model building portion. This could help us limit the total number of features used or it could prove that the more features the better the result.

```
In [21]: predictors15 = training[['basket_icon_click', 'basket_add_list', 'basket_add_detail', 'checked_delivery_detail', 'sign_in', 'saw_checkout', 'saw_sizecharts', 'device_mobile', 'device_tablet', 'device_computer', 'returning_user', 'loc_uk', 'ordered']]
```

```
In [22]: predictors02 = training.drop(['saw_sizecharts', 'device_mobile', 'device_tablet', 'device_computer', 'returning_user', 'loc_uk', 'ordered'])
```

```
In [23]: predictors15.head()
```

```
Out[23]:
```

	basket_icon_click	basket_add_list	basket_add_detail	list_size_dropdown	checked_delivery_detail
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	1	0	0	1

```
In [24]: predictors02.head()
```

```
Out[24]:
```

	basket_icon_click	basket_add_list	basket_add_detail	sort_by	image_picker	account_page_detail
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	1	0	1	0	0

```
In [25]: target = training['ordered']
```

```
In [26]: X_train15, X_test15, y_train15, y_test15 = train_test_split(predictors15, target
print( "Predictor - Training : ", X_train15.shape, "Predictor - Testing : ", X_t

Predictor - Training :  (341550, 8) Predictor - Testing :  (113851, 8)
```

```
In [27]: X_train02, X_test02, y_train02, y_test02 = train_test_split(predictors02, target
print( "Predictor - Training : ", X_train02.shape, "Predictor - Testing : ", X_t

Predictor - Training :  (341550, 20) Predictor - Testing :  (113851, 20)
```

Building a Predictions Model

```
In [28]: from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

```
In [29]: classifier = GaussianNB()
classifier = classifier.fit(X_train15, y_train15)
```

```
In [30]: predictions15 = classifier.predict(X_test15)
```

```
In [31]: confusion_matrix(y_test15, predictions15)
```


```
Out[31]: array([[108120,    984],
               [    61,   4686]])
```

```
In [32]: cm=confusion_matrix(y_test15, predictions15)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Gaussian Naive Bayes Confusion Matrix Over 0.15 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'B

plt.show()
```



```
In [33]: accuracy_score(y_test15, predictions15)
```

```
Out[33]: 0.9908213366593179
```

```
In [34]: classifier = GaussianNB()
classifier = classifier.fit(X_train02, y_train02)
```

```
In [35]: predictions02 = classifier.predict(X_test02)
```

```
In [36]: confusion_matrix(y_test02, predictions02)
```


```
Out[36]: array([[107744,   1306],
               [    55,   4746]])
```

```
In [37]: cm=confusion_matrix(y_test02, predictions02)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Gaussian Naive Bayes Confusion Matrix Over 0.02 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'Buy'])

plt.show()
```



```
In [38]: accuracy_score(y_test02, predictions02)
```

```
Out[38]: 0.9880457791323748
```

```
In [39]: log = LogisticRegression()
```

```
In [40]: log = log.fit(X_train15, y_train15)
```

```
In [41]: logPredict15 = log.predict(X_test15)
```

```
In [42]: cm=confusion_matrix(y_test15, logPredict15)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Logistic Regression Confusion Matrix Over 0.15 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'B

plt.show()
```

Logistic Regression Confusion Matrix Over 0.15 Correlation



```
In [43]: confusion_matrix(y_test15, logPredict15)
```

```
Out[43]: array([[108312,    792],
               [    47,   4700]])
```

```
In [44]: accuracy_score(y_test15, logPredict15)
```

```
Out[44]: 0.9926307190977681
```

```
In [45]: log = LogisticRegression()
log = log.fit(X_test02, y_test02)
```

```
In [46]: logPredict02 = log.predict(X_test02)
```


```
In [47]: cm=confusion_matrix(y_test02, logPredict02)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Logistic Regression Confusion Matrix Over 0.02 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'B

plt.show()
```

Logistic Regression Confusion Matrix Over 0.02 Correlation



```
In [48]: confusion_matrix(y_test02, logPredict02)
```

```
Out[48]: array([[108334,    716],
               [    66,   4735]])
```

```
In [49]: accuracy_score(y_test02, logPredict02)
```

```
Out[49]: 0.9931313734618054
```

```
In [50]: from sklearn.model_selection import GridSearchCV

LRparams = [{'penalty': ['none', 'l1', 'l2', 'elasticnet'], 'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000, 100000, 1000000]}
```

```
In [51]: LR_grid_search = GridSearchCV(estimator = LogisticRegression(),
                                       param_grid = LRparams,
                                       scoring = 'accuracy',
                                       cv = 10,
                                       n_jobs = -1)

LR_grid_search.fit(X_train15, y_train15)
bestLR = LR_grid_search.best_score_
bestParams = LR_grid_search.best_params_
print('Best Accuracy of Over 0.15 Correlation on Logistic Regression: ', bestLR)
print('Best Params of Over 0.15 Correlation on Logistic Regression: ', bestParam
```

```
/Users/wrasmussen/opt/anaconda3/lib/python3.7/site-packages/sklearn/model_selection/_search.py:925: UserWarning: One or more of the test scores are non-finite:
[0.99311667      nan 0.99195725      nan 0.99311667      nan
 0.99295564      nan 0.99311667      nan 0.99310789      nan
 0.99311667      nan 0.99311667      nan 0.99311667      nan
 0.99311667      nan 0.99311667      nan 0.99311667      nan]
category=UserWarning

/Users/wrasmussen/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:1323: UserWarning: Setting penalty='none' will ignore the C and l1_ratio parameters
"Setting penalty='none' will ignore the C and l1_ratio "
Best Accuracy of Over 0.15 Correlation on Logistic Regression: 0.99311667398623
92
```

Best Params of Over 0.15 Correlation on Logistic Regression: {'C': 0.001, 'penalty': 'none'}

```
In [52]: LR_grid_search = GridSearchCV(estimator = LogisticRegression(),
                                     param_grid = LRparams,
                                     scoring = 'accuracy',
                                     cv = 10,
                                     n_jobs = -1)

LR_grid_search.fit(X_train02, y_train02)
bestLR = LR_grid_search.best_score_
bestParams = LR_grid_search.best_params_
print('Best Accuracy of Over 0.02 Correlation on Logistic Regression: ', bestLR)
print('Best Params of Over 0.02 Correlation on Logistic Regression: ', bestParam
```

```
/Users/wrasmussen/opt/anaconda3/lib/python3.7/site-packages/sklearn/model_selection/_search.py:925: UserWarning: One or more of the test scores are non-finite:
[0.99315766      nan 0.99197775      nan 0.99315766      nan
 0.99297321      nan 0.99315766      nan 0.99313717      nan
 0.99315766      nan 0.99316059      nan 0.99315766      nan
 0.99315766      nan 0.99315766      nan 0.99315766      nan]
  category=UserWarning
Best Accuracy of Over 0.02 Correlation on Logistic Regression: 0.99316059142146
1
Best Params of Over 0.02 Correlation on Logistic Regression: {'C': 1, 'penalt
y': 'l2'}
```

```
In [53]: GNB_params = {'var_smoothing': np.logspace(0,-9, num=100)}
```

```
In [54]: GNB_grid_search = GridSearchCV(estimator = GaussianNB(),
                                     param_grid = GNB_params,
                                     scoring = 'accuracy',
                                     cv = 10,
                                     n_jobs = -1)

GNB_grid_search.fit(X_train15, y_train15)
bestGNB = GNB_grid_search.best_score_
bestParams = GNB_grid_search.best_params_
print('Best Accuracy of Over 0.15 Correlation on Logistic Regression: ', bestGNB)
print('Best Params of Over 0.15 Correlation on Logistic Regression: ', bestParam
```

```
Best Accuracy of Over 0.15 Correlation on Logistic Regression: 0.99292636510027
81
Best Params of Over 0.15 Correlation on Logistic Regression: {'var_smoothing':
0.008111308307896872}
```

```
In [55]: GNB_grid_search.fit(X_train02, y_train02)
bestGNB = GNB_grid_search.best_score_
bestParams = GNB_grid_search.best_params_
print('Best Accuracy of Over 0.02 Correlation on Logistic Regression: ', bestGNB)
print('Best Params of Over 0.02 Correlation on Logistic Regression: ', bestParam
```

```
Best Accuracy of Over 0.02 Correlation on Logistic Regression: 0.98857561118430
69
Best Params of Over 0.02 Correlation on Logistic Regression: {'var_smoothing':
0.001}
```

Based on the results of the hyperparameter tuning used here we can conclude that the Logistic

Regression we used is the best model for that. We can also see that for the Gaussian Naive Bayes we can increase our scores by using a var_smoothing of 0.008111308307896872. Let's run that model for use in propensity modeling to compare to Logistic Regression.

```
In [56]: gnb = GaussianNB(var_smoothing = 0.008111308307896872)
gnb = classifier.fit(X_train15, y_train15)
```

```
In [57]: predictions15 = gnb.predict(X_test15)
```

```
In [58]: confusion_matrix(y_test15, predictions15)
```


```
Out[58]: array([[108120,    984],
               [    61,   4686]])
```

```
In [59]: cm=confusion_matrix(y_test15, predictions15)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Gaussian Naive Bayes Confusion Matrix Over 0.15 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'B

plt.show()
```



```
In [60]: accuracy_score(y_test15, predictions15)
```

```
Out[60]: 0.9908213366593179
```

```
In [61]: log15 = LogisticRegression()
```

```
In [62]: log15 = log15.fit(X_train15, y_train15)
```

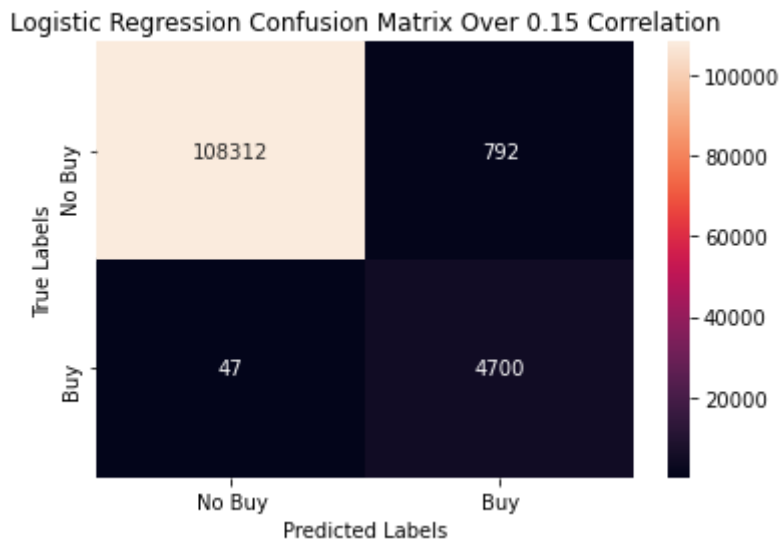
```
In [63]: logPredict15 = log15.predict(X_test15)
```

```
In [64]: cm=confusion_matrix(y_test15, logPredict15)

ax = plt.subplot()
sns.heatmap(cm, annot=True, fmt='g', ax=ax)

ax.set_xlabel('Predicted Labels');ax.set_ylabel('True Labels');
ax.set_title('Logistic Regression Confusion Matrix Over 0.15 Correlation');
ax.xaxis.set_ticklabels(['No Buy', 'Buy']);ax.yaxis.set_ticklabels(['No Buy', 'B

plt.show()
```



```
In [65]: accuracy_score(y_test15, logPredict15)
```

```
Out[65]: 0.9926307190977681
```

Based on the results of these models we can deduct that the Logistic Regression or Gaussian Naive Bayes with Over 0.15 Correlation are the strongest models. These will be used going forward in the propensity modeling.

Propensity Modeling

In this step we will be validating across a different data set from yesterday's shoppers to see the probability they would purchase.

```
In [66]: user_id = testing.UserID

yesterday = testing[['basket_icon_click', 'basket_add_list', 'basket_add_detail',
                    'checked delivery detail', 'sign in', 'saw checkout', 's
```

```
In [67]: yesterday.shape
```

```
Out[67]: (151655, 8)
```

```
In [68]: yesterday['propensity'] = log15.predict_proba(yesterday)[: ,1]
```

/Users/wrasmussen/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 """Entry point for launching an IPython kernel.

```
In [69]: yesterday.head()
```

```
Out[69]:
```

	basket_icon_click	basket_add_list	basket_add_detail	list_size_dropdown	checked_delivery_d
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	1	0	
4	0	0	0	0	

```
In [70]: yesterday.propensity.describe()
```

```
Out[70]:
```

count	1.516550e+05
mean	7.176702e-03
std	7.534979e-02
min	9.780716e-07
25%	3.342061e-06
50%	3.342061e-06
75%	4.299185e-06
max	9.598575e-01

Name: propensity, dtype: float64

```
In [106]: yesterday = pd.concat([user_id, yesterday], axis=1)
```

```
In [107]: yesterday.head()
```

```
Out[107]:
```

	UserID	basket_icon_click	basket_add_list	basket_add_detail	list_size_dropdown	checked_delivery_d
0	9d24-25k4-47889d24-25k4-494b-398124	0	0	0	0	
1	7732-1k58-47887732-1k58-4475-679678	0	0	0	0	


```
max          1.000000
Name: propensity, dtype: float64
```

```
In [108... yesterdayGNB = pd.concat([user_id, yesterdayGNB], axis=1)
```

```
In [109... target = yesterday[yesterday['propensity'] >= 0.5]
```

```
In [110... target.shape
```

Out[110... (1182, 10)

```
In [113... target.head()
```

Out[113...

	UserID	basket_icon_click	basket_add_list	basket_add_detail	list_size_dropdown	chec
	7j3d-j382-47157j3d-j382-4d3b-955343	1	0	1	1	
5						
	743b-08d2-4717743b-08d2-4634-230774	1	1	1	1	
23						
	b488-015d-472bb488-015d-4k88-211609	1	0	1	0	
58						
	7281-j047-47557281-j047-4425-872188	1	0	1	0	
162						
	0660-49k5-47890660-49k5-4070-438513	1	1	1	0	
287						

```
In [111... target1 = yesterdayGNB[yesterdayGNB['propensity'] >= 0.5]
```

```
In [112... target1.shape
```

Out[112... (1458, 10)

In [114...target1.head()

Out[114...

	UserID	basket_icon_click	basket_add_list	basket_add_detail	list_size_dropdown	chec
5	7j3d-j382-47157j3d-j382-4d3b-955343	1	0	1	1	
23	743b-08d2-4717743b-08d2-4634-230774	1	1	1	1	
58	b488-015d-472bb488-015d-4k88-211609	1	0	1	0	
162	7281-j047-47557281-j047-4425-872188	1	0	1	0	
287	0660-49k5-47890660-49k5-4070-438513	1	1	1	0	

Results

Logistic Regression

Logistic Regression identified 1182 different users from yesterday with a propensity to purchase over 0.5 which means they are statistically more likely to purchase than not purchase.

Gaussian Naive Bayes

Using Gaussian Naive Bayes found 1458 different users from yesterday with a propensity to purchase over 0.5 which means they are statistically more likely to purchase than not purchase.

Using These Results

From here now that we have the user_ids connected to the propensity to buy within the dataset we can start to target those users since they are most likely to purchase. We could export the dataframe in a csv format to pass onto sales to target or could export in another format as well.

In []: